

CS 415P/515 Parallel Programming

Assignment 2 Info

Jingke Li


Portland State University

Spring 2022

Prime-Finding Algorithm

Find all primes within a given bound n .

0. Consider all numbers from 2 to n as candidates.
1. Start with the first prime, 2, mark off all its multiples.
2. The next unmarked number is a new prime. The marking process continues for all primes up to \sqrt{n} (these are called sieve primes).
3. At the end, all remaining unmarked numbers are primes.



Sequential Version

```
prime.cpp
int main(int argc, char **argv) {
    int N = 100;
    cout << "Program prime over [2.." << N << "]" starting ...\n";

    bool candidate[N+1];
    for (int i = 2; i <= N; i++)
        candidate[i] = true;

    for (int i = 2; i <= sqrt(N); i++)
        if (candidate[i])
            for (int j = i+i; j <= N; j += i)
                candidate[j] = false;

    int totalPrimes = 0;
    for (int i = 2; i <= N; i++)
        if (candidate[i])
            totalPrimes++;
    cout << "Found " << totalPrimes << " primes\n";
}
```

Question: Which loops are parallelizable?

Parallelization Strategies

Consider the main loop-nest in the sequential program:

```
for (int i = 2; i <= sqrt(N); i++)  
    if (candidate[i])  
        for (int j = i+i; j <= N; j += i)  
            candidate[j] = false;
```

► **Strategy 1. Parallelizing the inner loop**

- natural parallelism, balanced workload, easy for OpenMP, ... *but*
- domain keeps changing, not as easy for explicit thread programming

► **Strategy 2. Parallelizing the outer loop**

- also easy for OpenMP, seems to work, ... *but*
- performs useless work, e.g. will assign threads to 4, 6, ...

Strategy 1 Improvement


- ▶ Have master to perform the task of finding sieves in $[2.. \sqrt{N}]$
- ▶ Partition a fixed domain $[\sqrt{N}..N]$ for the workers
- ▶ All threads can run concurrently

```
// Master: keep finding and broadcasting new sieve primes
for (int i = 2; i <= sqrt(N); i++)
    if (candidate[i])
        for (int j = i+i; j <= sqrt(N); j += i)
            candidate[j] = false;

// Workers: each works on a fixed section of (sqrt(N)..N)
for <each sieve prime p>
    for (int j = start; j <= end; j += p)
        candidate[j] = false;
```

A Closer Look at Strategy 2

Assign a thread to each iteration of the outer loop



- All threads work independently; no synchronization is needed
- Threads will work on composite numbers, wasting resource

Question: Is there a race condition?

(Multiple threads may mark the same cell.)

Strategy 2 Improvement

Assign threads to primes only

- ▶ Start a single thread first (for sieving prime 2)
- ▶ As soon as a new sieve prime is found, a new thread may start

